

CIND820_project

December 1, 2021

INITIALIZE - import Python modules *****

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn import linear_model
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge, RidgeCV
```

```
[2]: pip install xgboost
```

Requirement already satisfied: xgboost in /opt/conda/lib/python3.7/site-packages (1.5.1)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.4.1)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.18.4)

Note: you may need to restart the kernel to use updated packages.

```
[3]: import xgboost
from xgboost import XGBRegressor
```

IMPORT DATA - import data from csv files - combine into 2 dataframes: -
Regular season - Postseason - check summary statistics for resulting dataframes

```
[4]: #Import data

df_RS_sum = pd.read_csv('Summary_RS.csv', encoding='latin-1')
df_RS_pen = pd.read_csv('Penalties_RS.csv', encoding='latin-1')
df_RS_tgg = pd.read_csv('Team Goal-Games_RS.csv', encoding='latin-1')
df_RS_satc = pd.read_csv('SAT Counts_RS.csv', encoding='latin-1')
df_RS_satp = pd.read_csv('SAT Percentages_RS.csv', encoding='latin-1')

df_PS_sum = pd.read_csv('Summary_PS.csv', encoding='latin-1')
```

```
df_PS_pen = pd.read_csv('Penalties_PS.csv', encoding='latin-1')
df_PS_tgg = pd.read_csv('Team Goal-Games_PS.csv', encoding='latin-1')
df_PS_satc = pd.read_csv('SAT Counts_PS.csv', encoding='latin-1')
df_PS_satp = pd.read_csv('SAT Percentages_PS.csv', encoding='latin-1')
```

```
[5]: #Check import sample
df_PS_sum.head()
```

```
[5]:
```

	Team	Season	GP	W	L	T	OT	P	P%	RW	...	\
0	Tampa Bay Lightning	20202021	23	16	7	--	--	32.0	0.696	16	...	
1	Montréal Canadiens	20202021	22	13	9	--	--	26.0	0.591	7	...	
2	New York Islanders	20202021	19	11	8	--	--	22.0	0.579	7	...	
3	Vegas Golden Knights	20202021	19	10	9	--	--	20.0	0.526	8	...	
4	Boston Bruins	20202021	11	6	5	--	--	12.0	0.545	3	...	

	GA	GF/GP	GA/GP	PP%	PK%	Net PP%	Net PK%	Shots/GP	SA/GP	FOW%
0	45.0	3.26	1.96	32.4	84.1	32.4	87.3	29.6	30.4	48.1
1	54.0	2.32	2.45	18.9	91.8	17.0	98.4	28.9	30.7	49.4
2	53.0	2.84	2.79	20.4	65.1	16.3	65.1	28.5	35.2	51.8
3	46.0	2.79	2.42	9.3	71.8	9.3	74.4	32.1	26.4	53.5
4	32.0	3.00	2.91	36.4	75.7	36.4	75.7	35.8	31.0	52.3

[5 rows x 23 columns]

```
[6]: #Create regular season dataframe

#Rename column with shared name but different data
df_RS_satp = df_RS_satp.rename({'GF':'5v5 GF', 'GA':'5v5 GA'}, axis=1)

#Drop duplicate columns
df_RS_pen = df_RS_pen.drop(['GP','W','L','T','OT','P'], axis=1)
df_RS_satc = df_RS_satc.drop(['GP'], axis=1)
df_RS_satp = df_RS_satp.drop(['GP','P','P%'], axis=1)
df_RS_tgg = df_RS_tgg.drop(['GP','W','L','T','OT','P','P%'], axis=1)

#Merge dataframes to create one regular season dataframe
df = pd.merge(df_RS_sum, df_RS_pen, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_satc, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_satp, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_tgg, how='inner', on=['Team','Season'])

#Drop irrelevant columns (Ties not applicable in seasons being studied)
df = df.drop(['T'], axis=1)

#Convert season column to string
df['Season'] = df['Season'].apply(str)
```

```
#Convert SAT for/against to rate per game
df['SAT For/GP'] = df['SAT For'] / df['GP']
df['SAT Agst/GP'] = df['SAT Agst'] / df['GP']
```

```
[7]: #Check summary statistics for regular season dataframe
df.describe()
```

```
[7]:
```

	GP	W	L	OT	P	P%	\
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000	
mean	74.311688	37.155844	28.707792	8.448052	82.759740	0.556929	
std	10.380676	8.935594	7.828994	2.906028	17.944379	0.094176	
min	56.000000	15.000000	12.000000	2.000000	37.000000	0.275000	
25%	69.000000	31.000000	24.000000	7.000000	72.000000	0.488000	
50%	82.000000	37.000000	28.000000	8.000000	81.500000	0.569000	
75%	82.000000	44.000000	34.000000	10.000000	97.750000	0.622750	
max	82.000000	62.000000	56.000000	15.000000	128.000000	0.780000	

	RW	ROW	S/O Win	GF	...	Win% 3 Goal Game	\
count	154.000000	154.000000	154.000000	154.000000	...	154.000000	
mean	28.707792	34.298701	2.857143	215.642857	...	0.499370	
std	8.100613	8.792547	1.693620	38.727923	...	0.163777	
min	11.000000	11.000000	0.000000	124.000000	...	0.032000	
25%	23.000000	28.000000	2.000000	189.250000	...	0.385500	
50%	29.000000	34.000000	3.000000	221.000000	...	0.515500	
75%	35.000000	41.000000	4.000000	243.000000	...	0.618000	
max	49.000000	56.000000	9.000000	319.000000	...	0.875000	

	Wins 1 Goal Game	Wins 2 Goal Game	Wins 3 Goal Game	Loss 1 Goal Game	\
count	154.000000	154.000000	154.000000	154.000000	
mean	16.012987	7.733766	13.409091	7.564935	
std	4.216350	2.985898	5.188743	2.976872	
min	5.000000	0.000000	1.000000	1.000000	
25%	13.000000	6.000000	10.000000	5.000000	
50%	16.000000	7.000000	14.000000	8.000000	
75%	19.000000	10.000000	17.000000	9.000000	
max	25.000000	16.000000	30.000000	15.000000	

	Loss 2 Goal Game	Loss 3 Goal Game	OT Loss 1 Goal Game	SAT For/GP	\
count	154.000000	154.000000	154.000000	154.000000	
mean	7.733766	13.409091	8.448052	44.869488	
std	3.031517	5.099806	2.906028	3.125934	
min	2.000000	2.000000	2.000000	37.321429	
25%	5.250000	10.000000	7.000000	42.664634	
50%	8.000000	13.000000	8.000000	44.622570	
75%	10.000000	16.000000	10.000000	46.908537	
max	16.000000	30.000000	15.000000	53.609756	

	SAT Agst/GP
count	154.000000
mean	44.867717
std	3.010061
min	35.053571
25%	42.871951
50%	45.009364
75%	47.088415
max	53.914634

[8 rows x 75 columns]

[8]: *#Create postseason dataframe following same process as regular season*

```
df_PS_satp = df_PS_satp.rename({'GF': '5v5 GF', 'GA': '5v5 GA'}, axis=1)

df_PS_pen = df_PS_pen.drop(['GP', 'W', 'L', 'T', 'OT', 'P'], axis=1)
df_PS_satc = df_PS_satc.drop(['GP'], axis=1)
df_PS_satp = df_PS_satp.drop(['GP', 'P', 'P%'], axis=1)
df_PS_tgg = df_PS_tgg.drop(['GP', 'W', 'L', 'T', 'P', 'P%'], axis=1)

df2 = pd.merge(df_PS_sum, df_PS_pen, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_satc, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_satp, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_tgg, how='inner', on=['Team', 'Season'])

df2 = df2.drop(['T'], axis=1)

df2['Season'] = df2['Season'].apply(str)

df2['SAT For/GP'] = df2['SAT For'] / df2['GP']
df2['SAT Agst/GP'] = df2['SAT Agst'] / df2['GP']
```

[9]: df2.describe()

[9]:

	GP	W	L	P	P%	RW \
count	88.000000	88.000000	88.000000	88.000000	88.000000	88.000000
mean	10.727273	5.363636	5.329545	5.117261	2.505977	4.738636
std	6.629310	4.849889	1.981069	8.387413	3.631763	4.450221
min	3.000000	0.000000	3.000000	0.000000	0.000000	0.000000
25%	5.750000	2.000000	4.000000	0.333000	0.333000	1.000000
50%	8.500000	4.000000	4.000000	0.607500	0.612500	3.000000
75%	15.000000	8.250000	6.250000	6.500000	3.250000	7.000000
max	27.000000	18.000000	12.000000	36.000000	14.000000	16.000000

	ROW	S/O Win	GF	GA	...	5v5 S%+Sv% \
count	88.000000	88.000000	88.000000	88.000000	...	88.000000

mean	2.409091	16.272727	29.534091	14.671023	...	98.932955
std	4.173359	21.981041	18.536047	17.144004	...	3.018781
min	0.000000	0.000000	4.000000	0.750000	...	90.600000
25%	0.000000	0.000000	16.000000	2.652500	...	97.150000
50%	0.000000	7.000000	23.000000	3.355000	...	99.850000
75%	3.250000	26.250000	42.000000	22.250000	...	101.000000
max	17.000000	86.000000	77.000000	82.000000	...	103.800000

	Wins 2 Goal Game	Wins 3 Goal Game	Loss 1 Goal Game	Loss 2 Goal Game \
count	88.000000	88.000000	88.000000	88.000000
mean	1.829545	1.397727	2.136364	1.795455
std	2.354851	1.847893	1.709889	1.576837
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000	1.000000
50%	1.000000	1.000000	2.000000	1.500000
75%	2.000000	2.000000	3.000000	3.000000
max	12.000000	9.000000	6.000000	7.000000

	Loss 3 Goal Game	OT Loss 1 Goal Game	Unnamed: 18	SAT For/GP \
count	88.000000	88.000000	40.000000	88.000000
mean	1.397727	0.806818	0.075000	47.657206
std	1.255240	1.239744	0.266747	6.483257
min	0.000000	0.000000	0.000000	36.571429
25%	0.000000	0.000000	0.000000	43.221429
50%	1.000000	0.000000	0.000000	46.900219
75%	2.000000	1.250000	0.000000	51.723214
max	5.000000	5.000000	1.000000	69.500000

	SAT Agst/GP
count	88.000000
mean	47.749411
std	6.219736
min	36.400000
25%	44.250000
50%	47.431373
75%	50.987500
max	66.900000

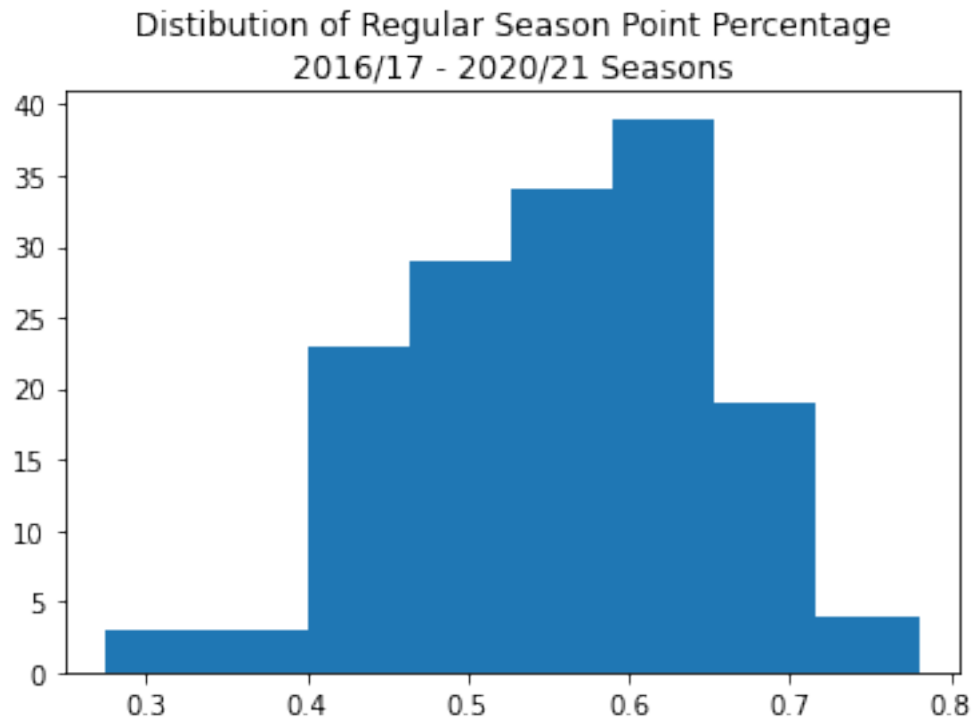
[8 rows x 72 columns]

```
[10]: fig, ax = plt.subplots()

plt.hist(df['P%'], bins=8)

plt.title('Distribution of Regular Season Point Percentage\n2016/17 - 2020/21_\n↳Seasons')

plt.savefig('hist_RS_Ppct.png', dpi=144)
```

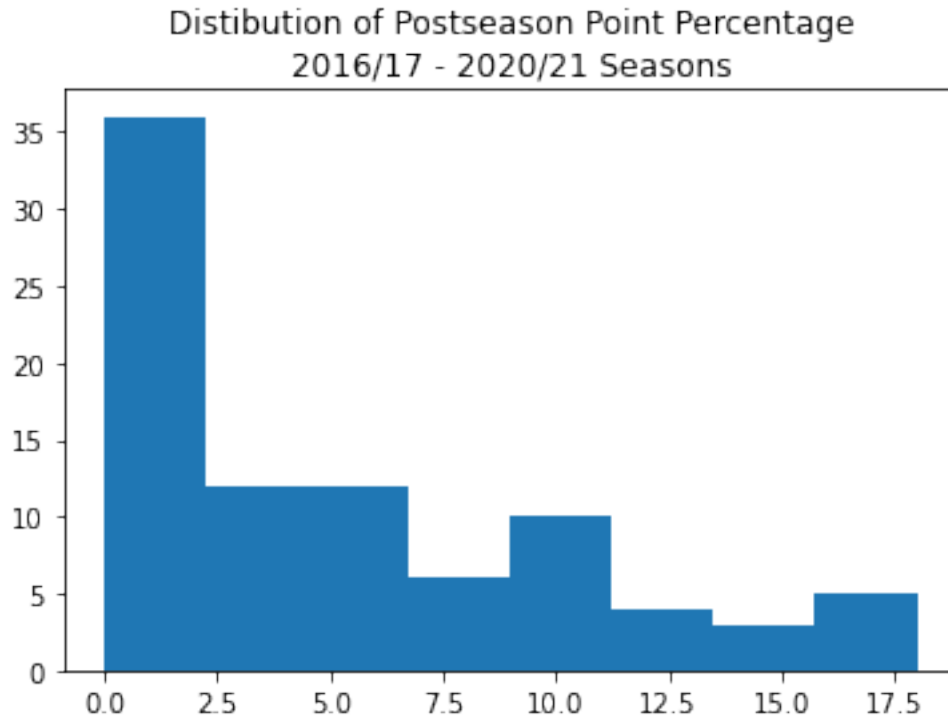


```
[11]: fig, ax = plt.subplots()

plt.hist(df2['W'], bins=8)

plt.title('Distribution of Postseason Point Percentage\n2016/17 - 2020/21_\n↪Seasons')

plt.savefig('hist_PS_W.png', dpi=144)
```

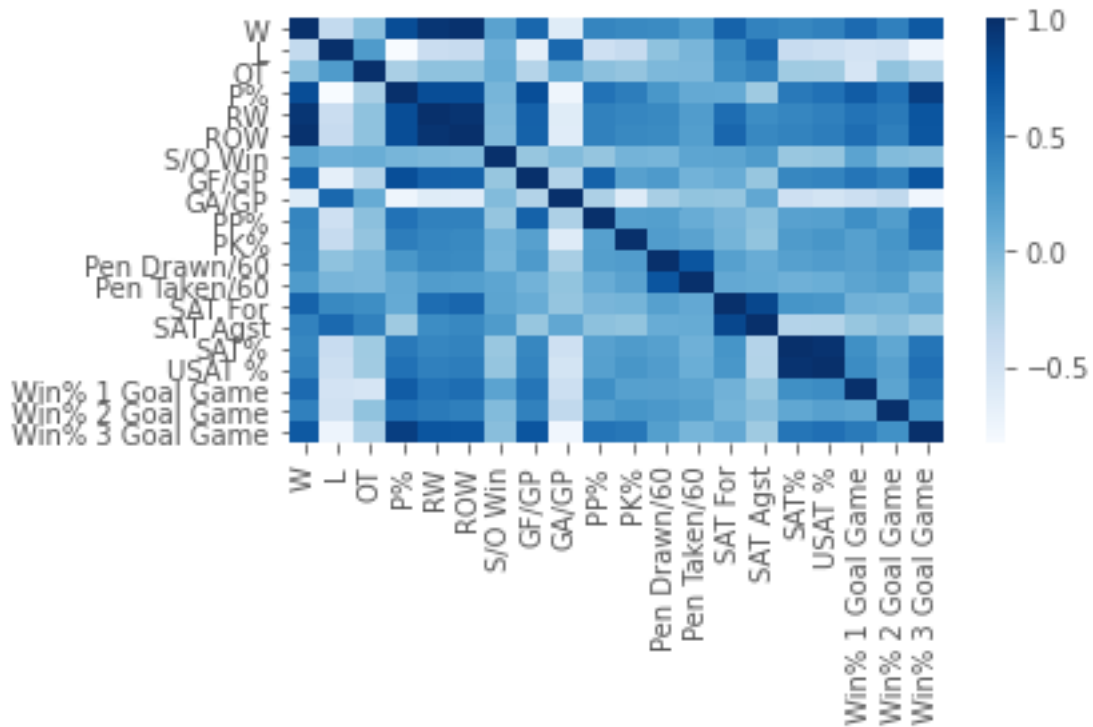


INITIAL DATA ASSESSMENT - reduce dataframes to set of variables expected to be of primary interest. Preferentially select variables as rates or ratios from each category of statistics. - create heat map to check correlation between selected variables *****

```
[12]: #Reduce regular season dataframe for initial heatmap
df_reg = df[['W','L','OT','P%', 'RW','ROW','S/O Win','GF/GP','GA/
↳GP','PP%','PK%','Pen Drawn/60','Pen Taken/60','SAT For','SAT_
↳Agst','SAT%','USAT %','Win% 1 Goal Game','Win% 2 Goal Game', 'Win% 3 Goal_
↳Game']]
```

```
[13]: #Create correlation heat map for regular season
import matplotlib
matplotlib.use('Agg')
import matplotlib.pyplot as plt
matplotlib.style.use('ggplot')
import seaborn as sns
corr = df_reg.corr()
sns_plot = sns.heatmap(corr, cmap="Blues", annot=False)

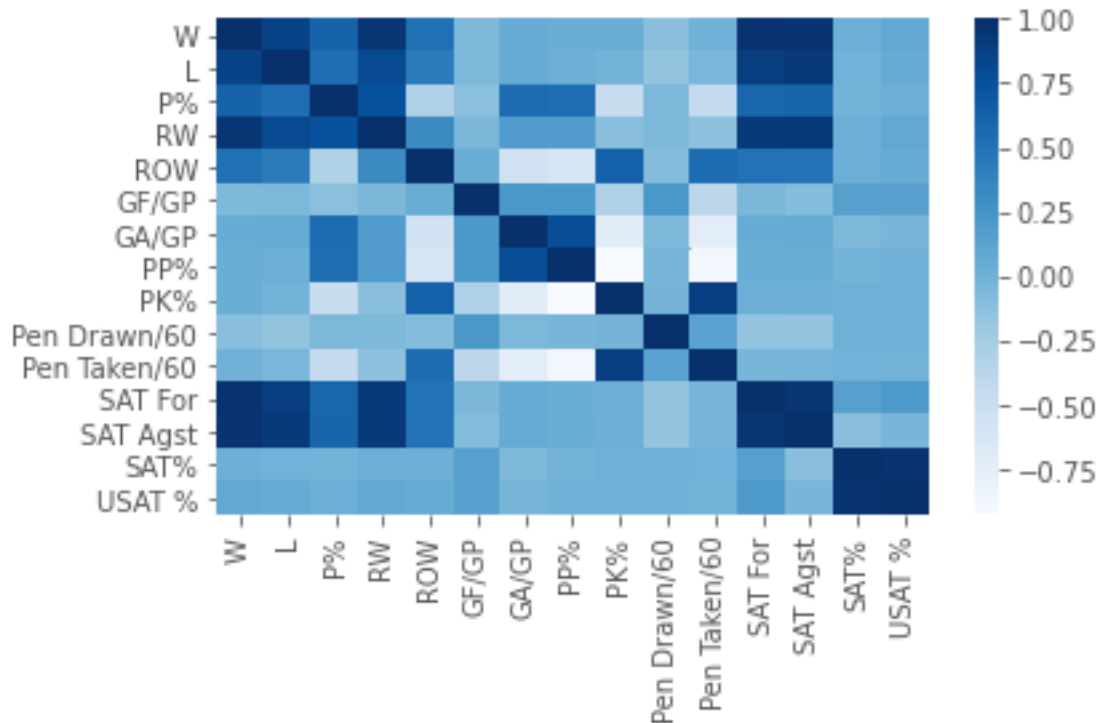
fig = sns_plot.get_figure()
fig.tight_layout()
fig.savefig("reg_heat.png")
```



```
[14]: #Reduce postseason dataframe for initial heatmap
df_post = df2[['W', 'L', 'OT', 'P%', 'RW', 'ROW', 'GF/GP', 'GA/GP', 'PP%', 'PK%', 'Pen_
↳ Drawn/60', 'Pen Taken/60', 'SAT For', 'SAT Agst', 'SAT%', 'USAT %', 'Win% 1 Goal_
↳ Game', 'Win% 2 Goal Game', 'Win% 3 Goal Game']]
```

```
[15]: #Create correlation heat map for postseason
corr = df_post.corr()
sns_plot = sns.heatmap(corr, cmap="Blues", annot=False)

fig = sns_plot.get_figure()
fig.tight_layout()
fig.savefig("post_heat.png")
```

MODELLING: REGULAR SEASON Initial model: - Train initial model for regular season using goal-based input variables - Test multiple types of regression models to find best fit. Considering linear regression, XGBoost regression, ridge regression Additional models: - Add additional variables to initial model in categorical groups to test for improved models.

- Additional categorical groups include shot metrics, penalties, special teams, records in close games - Additional models without using goal metrics

```
[16]: #TRAIN INITIAL MODEL MO
#Linear Regression
#Inputs: goal based variables
#Target: Regular season P%
#70/30 Train/test split

#Target Variable = Roint Percentage
reg_col_name = 'P%'

#Select feature variables
feature_names = ['GF/GP', 'GA/GP']

#Use 70/30 train/test split
```

```

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳df[reg_col_name], test_size=0.3, random_state=2)

#Train model & predict test set
MO = linear_model.LinearRegression()
MO.fit(X_train, y_train)

y_pred = MO.predict(X_test)

```

```

[17]: #Check model metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
↳y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.019301933879415853
Mean Squared Error: 0.0005326454997485497
Root Mean Squared Error: 0.02307911392901707
R-Squared: 0.9331582772976956
Adjusted R-Squared: 0.9318728595534206

```

[18]: #TRAIN INITIAL MODEL MO with kfold cross validation
#Linear Regression
#Inputs: goal based variables
#Target: Regular season P%
#kfold cross validation

X, y = df[['GF/GP', 'GA/GP']], df['P%']

#Train model & predict test set
MOK = linear_model.LinearRegression()
#MOK.fit(X, y)

#y_pred = MOK.predict(X)

from sklearn.model_selection import RepeatedKFold

cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

from sklearn.model_selection import cross_val_score

scores = cross_val_score(MOK, X, y, scoring='neg_mean_absolute_error', cv=cv)
scores.mean()

```

[18]: -0.019066608315764132

```
[19]: #MODEL MO_XGB
#XGBoost Regression
#Inputs: goal based variables
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

MO_XGB = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
    max_depth = 4, n_estimators = 100)

MO_XGB.fit(X_train, y_train)

y_pred = MO_XGB.predict(X_test)
```

```
[20]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.024832683020449716
Mean Squared Error: 0.0010186796962558568
Root Mean Squared Error: 0.031916761995162614
R-Squared: 0.8721658104466374
Adjusted R-Squared: 0.8697074606475343

```
[21]: #MODEL MORR
#Ridge Regression
#Inputs: goal based variables
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)
```

```
MORR = Ridge(alpha=0.1)
MORR.fit(X_train,y_train)

y_pred = MORR.predict(X_test)
```

```
[22]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

```
Mean Absolute Error: 0.019349022473394435
Mean Squared Error: 0.0005375438606248156
Root Mean Squared Error: 0.02318499214200461
R-Squared: 0.93254358163324
Adjusted R-Squared: 0.9312463428184946
```

```
[23]: #Continue regular season prediction adding additional variables to goal-based
    ↪model. Linear Regression dropped due to suspected overfitting in M0.
```

```
[24]: #Does adding shot metric variables improve goal based model?
#MODEL M1
#XGBoost Regression
#Inputs: goal based variables + shot-based metrics
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M1 = XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
                    max_depth = 2, n_estimators = 100)
M1.fit(X_train,y_train)

y_pred = M1.predict(X_test)
```

```
[25]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
```

```
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
→y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.024218783079309674
Mean Squared Error: 0.0010224722293729596
Root Mean Squared Error: 0.03197611967348383
R-Squared: 0.8716898851885204
Adjusted R-Squared: 0.865337899306764

```
[26]: #MODEL M1RR
#Ridge Regression
#Inputs: goal based variables + shot-based metrics
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
→df[reg_col_name], test_size=0.3, random_state=2)

M1RR = Ridge(alpha=1)
M1RR.fit(X_train, y_train)

y_pred = M1RR.predict(X_test)
```

```
[27]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
→y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
→y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.020269714102022034
Mean Squared Error: 0.0006176650868208685
Root Mean Squared Error: 0.024852868784526035
R-Squared: 0.9224891631006638
Adjusted R-Squared: 0.9186519929571323

```
[28]: #Can shot metrics provide predictive ability without goal based data?
#MODEL M2
#XGBoost Regression
#Inputs: shot metrics
#Target: Regular season P%
```

```

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M2 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.3,
    max_depth = 5, n_estimators = 100)
M2.fit(X_train, y_train)

y_pred = M2.predict(X_test)

```

```

[29]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.06372146529847003
 Mean Squared Error: 0.0056991383192163516
 Root Mean Squared Error: 0.07549263751662377
 R-Squared: 0.28481471568806993
 Adjusted R-Squared: 0.26398407633917875

```

[30]: #MODEL M2RR
      #Ridge Regression
      #Inputs: shot metrics
      #Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M2RR = Ridge(alpha=0.1)
M2RR.fit(X_train, y_train)

y_pred = M2RR.predict(X_test)

```

```

[31]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))

```

```

print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.07074727860750397
 Mean Squared Error: 0.006334153385074077
 Root Mean Squared Error: 0.0795873946368021
 R-Squared: 0.20512662865103448
 Adjusted R-Squared: 0.1819749770583462

```

[32]: #Do S% and SV% improve shot based model?
      #MODEL M4
      #XGBoost Regression
      #Inputs: shot metrics, shooting/save percentage
      #Target: Regular season P%

      reg_col_name = 'P%'

      feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

      M4 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.3,
                        max_depth = 3, n_estimators = 10)
      M4.fit(X_train, y_train)

      y_pred = M4.predict(X_test)

```

```

[33]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.04410175598935877
 Mean Squared Error: 0.003107177336544678
 Root Mean Squared Error: 0.05574206074899526
 R-Squared: 0.6100800888879196
 Adjusted R-Squared: 0.590777122991282

```
[34]: #MODEL M4RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M4RR = Ridge(alpha=0.1)
M4RR.fit(X_train, y_train)

y_pred = M4RR.predict(X_test)
```

```
[35]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.03219329522690666
Mean Squared Error: 0.0017621542671799556
Root Mean Squared Error: 0.041978021239452865
R-Squared: 0.7788671321899294
Adjusted R-Squared: 0.7679199605161635

```
[36]: #Does adding special teams performance improve previous model?
#MODEL M5
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
    ↪'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M5 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.5,
    ↪learning_rate = 0.1,
```



```

        max_depth = 4, n_estimators = 100)
M5.fit(X_train,y_train)

y_pred = M5.predict(X_test)

```

```

[37]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
      ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.037449044826182915
 Mean Squared Error: 0.002268004998376496
 Root Mean Squared Error: 0.04762357607715422
 R-Squared: 0.7153878869520387
 Adjusted R-Squared: 0.6952637981506677

```

[38]: #MODEL M5RR
      #Ridge Regression
      #Inputs: shot metrics, shooting/save percentage, special teams
      #Target: Regular season P%

      reg_col_name = 'P%'

      feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
      ↪'PP%', 'PK%']

      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
      ↪df[reg_col_name], test_size=0.3,random_state=2)

      M5RR = Ridge(alpha=0.1)
      M5RR.fit(X_train,y_train)

      y_pred = M5RR.predict(X_test)

```

```

[39]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
      ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 0.02511585765061583
 Mean Squared Error: 0.0011778810039069912

Root Mean Squared Error: 0.03432027103487079
R-Squared: 0.8521876267111418
Adjusted R-Squared: 0.8417362467816265

```
[40]: #Does adding penalty rates improve the previous model?
#MODEL M6
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳ df[reg_col_name], test_size=0.3, random_state=2)

M6 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.5,
↳ learning_rate = 0.1,
                    max_depth = 40, n_estimators = 100)
M6.fit(X_train, y_train)

y_pred = M6.predict(X_test)
```

```
[41]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
↳ y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.04500114227355795
Mean Squared Error: 0.0030528095528431706
Root Mean Squared Error: 0.05525223572710131
R-Squared: 0.6169027060391592
Adjusted R-Squared: 0.5813575962902152

```
[42]: #MODEL M6RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Regular season P%

reg_col_name = 'P%'
```

```

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳ df[reg_col_name], test_size=0.3, random_state=2)

M6RR = Ridge(alpha=0.1)
M6RR.fit(X_train, y_train)

y_pred = M6RR.predict(X_test)

```

```

[43]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1 - metrics.r2_score(y_test,
↳ y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))

```

Mean Absolute Error: 0.025264003393739975
Mean Squared Error: 0.0011659666653443162
Root Mean Squared Error: 0.03414625404556576
R-Squared: 0.8536827579283653
Adjusted R-Squared: 0.8401069313443992

```

[44]: #Does kfold cross validation help model M5RR?
#MODEL M5RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams
#Target: Regular season P%

X, y = df[['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%']], df['P%']

#Train model & predict test set
M5RRk = Ridge(alpha=0.1)
#MOK.fit(X,y)

#y_pred = MOK.predict(X)

from sklearn.model_selection import RepeatedKFold

cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

from sklearn.model_selection import cross_val_score

scores = cross_val_score(M5RRk, X, y, scoring='neg_mean_absolute_error', cv=cv)

```

```
scores.mean()
```

```
[44]: -0.02964831601405144
```

POSTSEASON MODELS - Model postseason, using wins as target variable. - Use input variables matching regular season models: initial base model and other most successful regular season models
- Results to be compared with regular season models to assess impact of small postseason sample sizes *****

```
[45]: #Can the inputs for the most successful regular season model M5 provide similar
      ↪ results on postseason data? Use W instead of P% for target variable for
      ↪ postseason.
      #MODEL P5
      #XGBoost Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penalty rates
      #Target: Postseason W

      reg_col_name = 'W'

      feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
      ↪ 'PP%', 'PK%']

      X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
      ↪ df2[reg_col_name], test_size=0.3, random_state=2)

      P5 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
      ↪ learning_rate = 0.1,
      ↪ max_depth = 4, n_estimators = 100)
      P5.fit(X_train, y_train)

      y_pred = P5.predict(X_test)
```

```
[46]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪ y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
      ↪ y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

```
Mean Absolute Error: 2.841516375541687
Mean Squared Error: 15.753672860832477
Root Mean Squared Error: 3.9690896765924144
R-Squared: 0.272492872447303
Adjusted R-Squared: 0.17640702541204112
```

```
[47]: #MODEL P5RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
↳ df2[reg_col_name], test_size=0.3, random_state=2)

P5RR = Ridge(alpha=0.1)
P5RR.fit(X_train, y_train)

y_pred = P5RR.predict(X_test)
```

```
[48]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
↳ y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

```
Mean Absolute Error: 2.8316068910219565
Mean Squared Error: 12.136678677422992
Root Mean Squared Error: 3.4837736260301115
R-Squared: 0.43952624123645256
Adjusted R-Squared: 0.3655014051733426
```

```
[49]: #Does goal based model have similar reduction in results?
#MODEL P0
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
↳ df2[reg_col_name], test_size=0.3, random_state=2)

P0 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
↳ learning_rate = 0.1,
```

```

        max_depth = 4, n_estimators = 100)
P0.fit(X_train,y_train)

y_pred = P0.predict(X_test)

```

```

[50]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
      ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 3.720087496218858
 Mean Squared Error: 20.83882289016658
 Root Mean Squared Error: 4.564955957089463
 R-Squared: 0.03765983232412051
 Adjusted R-Squared: 0.004475688611159212

```

[51]: #MODEL PORR
      #Ridge Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
      #Target: Postseason W

      reg_col_name = 'W'

      feature_names = ['GF/GP', 'GA/GP']

      X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
      ↪df2[reg_col_name], test_size=0.3,random_state=2)

      PORR = Ridge(alpha=0.5)
      PORR.fit(X_train,y_train)

      y_pred = PORR.predict(X_test)

```

```

[52]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
      ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

Mean Absolute Error: 3.76161208247753
 Mean Squared Error: 23.979697311732068
 Root Mean Squared Error: 4.896906912708477

R-Squared: -0.10738624985763812
Adjusted R-Squared: -0.14557198261134996

PREDICT POSTSEASON SUCCESS USING REGULAR SEASON INPUT VARIABLES - Use regular season input variables to predict postseason success. Chose variables based on results of previous models. *****

```
[53]: #Use regular season statistics for input variables
RS = df[['Team', 'Season', 'RW', 'P', 'GF/GP', 'GA/GP', 'Win% 1 Goal Game',
        ↪ 'Win% 2 Goal Game', 'Win% 3 Goal Game', 'SAT For/GP', 'SAT Agst/GP', 'SAT%',
        ↪ '5v5 S%', '5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]

#Use postseason W as target variable
PS = df2[['Team', 'Season', 'P%', 'W']]

#Create new dataframe
df3 = pd.merge(RS, PS, on=['Team', 'Season'])
```

```
[54]: #Can regular season inputs be used to predict postseason W?
#MODEL RP5
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
        ↪ 'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
        ↪ df3[reg_col_name], test_size=0.3, random_state=2)

RP5 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.5,
        ↪ learning_rate = 0.1,
        max_depth = 2, n_estimators = 100)
RP5.fit(X_train, y_train)

y_pred = RP5.predict(X_test)
```

```
[55]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
        ↪ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
        ↪ y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 5.222498823095251
Mean Squared Error: 42.04227377577858
Root Mean Squared Error: 6.484001370741572
R-Squared: -0.5449550147465765
Adjusted R-Squared: -0.749005677071596

```
[56]: #MODEL RP5RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penalty rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
                 'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
                                                    df3[reg_col_name], test_size=0.3, random_state=2)

RP5RR = Ridge(alpha=0.1)
RP5RR.fit(X_train, y_train)

y_pred = RP5RR.predict(X_test)
```

```
[57]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
                               y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1 - metrics.r2_score(y_test,
                               y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))
```

Mean Absolute Error: 5.132641301308711
Mean Squared Error: 39.62557528535863
Root Mean Squared Error: 6.294884850841883
R-Squared: -0.45614700993176927
Adjusted R-Squared: -0.6484683131303048

```
[58]: #MODEL RPO
#XGBoost Regression
#Inputs: goal rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['GF/GP', 'GA/GP']
```



```

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
↳df3[reg_col_name], test_size=0.3,random_state=2)

RPO = XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.5,
↳learning_rate = 0.1,
                    max_depth = 10, n_estimators = 100)
RPO.fit(X_train,y_train)

y_pred = RPO.predict(X_test)

```

```

[59]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
↳y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

```

Mean Absolute Error: 5.422038421862656
Mean Squared Error: 48.07547664870346
Root Mean Squared Error: 6.933648148608599
R-Squared: -0.7666610785817534
Adjusted R-Squared: -0.8275804261190554

```

```

[60]: #MODEL RPORR
      #Ridge Regression
      #Inputs: goal rates
      #Target: Postseason W

      reg_col_name = 'W'

      feature_names = ['GF/GP', 'GA/GP']

      X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
↳df3[reg_col_name], test_size=0.3,random_state=2)

      RPO = Ridge(alpha=0.1)
      RPO.fit(X_train,y_train)

      y_pred = RPO.predict(X_test)

```

```

[61]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

```
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
→y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 4.693376720415959
Mean Squared Error: 35.838794498087864
Root Mean Squared Error: 5.986551135510985
R-Squared: -0.3169916921618132
Adjusted R-Squared: -0.3624051987880825

CHECK REGULAR SEASON TO POSTSEASON CORRELATIONS - check correlation of each variable used as input variable to previous model between regular season and postseason. - higher correlation may suggest higher predictive power in postseason prediction models.

```
[62]: #Check correlation between key variables: regular season vs postseason

#Combine data
reg_shot = df[['Team', 'Season', 'SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%',
→'5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]
post_shot = df2[['Team', 'Season', 'SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5
→S%', '5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]
post_shot = post_shot.rename(columns={'SAT For/GP': 'SAT For/GP P', 'SAT Agst/
→GP': 'SAT Agst/GP P', 'SAT%': 'SAT% P', '5v5 S%': '5v5 S% P', '5v5 Sv%': '5v5
→Sv% P', 'PP%': 'PP% P', 'PK%': 'PK% P', 'Pen Drawn/60': 'Pen Drawn/60 P', 'Pen
→Taken/60': 'Pen Taken/60 P'})
compare = pd.merge(reg_shot, post_shot, on=['Team', 'Season'])

#Check correlations
from scipy.stats import pearsonr
stats = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%', 'PP%', 'PK%',
→'Pen Drawn/60', 'Pen Taken/60']

for i in stats:
    list1 = compare[i]
    list2 = compare[i + ' P']

    corr, _ = pearsonr(list1, list2)
    print('Pearsons correlation (' + i + '): %.3f' % corr)
```

Pearsons correlation (SAT For/GP): 0.261
Pearsons correlation (SAT Agst/GP): 0.254
Pearsons correlation (SAT%): 0.466
Pearsons correlation (5v5 S%): -0.080
Pearsons correlation (5v5 Sv%): 0.043
Pearsons correlation (PP%): -0.063
Pearsons correlation (PK%): -0.106
Pearsons correlation (Pen Drawn/60): 0.393

Pearsons correlation (Pen Taken/60): -0.125

```
[63]: #Does model improve by removing variables with poor regular season to  
→postseason correlation?  
#MODEL RP5RR_simple  
#Ridge Regression  
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates  
#Target: Postseason W  
  
reg_col_name = 'W'  
  
feature_names = ['SAT%']  
  
X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],  
→df3[reg_col_name], test_size=0.3,random_state=2)  
  
RP5RR_simple = Ridge(alpha=0.1)  
RP5RR_simple.fit(X_train,y_train)  
  
y_pred = RP5RR_simple.predict(X_test)
```

```
[64]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,  
→y_pred)))  
print('R-Squared:', metrics.r2_score(y_test, y_pred))  
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,  
→y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 4.992943710259225
Mean Squared Error: 38.76544416640514
Root Mean Squared Error: 6.226190180712853
R-Squared: -0.4245392074457781
Adjusted R-Squared: -0.44868393977536747

```
[65]: #MODEL RP5RR_simple  
#Ridge Regression  
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates  
#Target: Postseason W  
  
reg_col_name = 'W'  
  
feature_names = ['SAT%']  
  
X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],  
→df3[reg_col_name], test_size=0.3,random_state=2)
```

```

RP_simple = XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.5,
    ↪learning_rate = 0.1,
    max_depth = 10, n_estimators = 100)
RP_simple.fit(X_train,y_train)

y_pred = RP_simple.predict(X_test)

```

```

[66]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
      print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
    ↪y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))

```

```

Mean Absolute Error: 6.059022011718264
Mean Squared Error: 56.448535744619136
Root Mean Squared Error: 7.513224057927404
R-Squared: -1.074351374020937
Adjusted R-Squared: -1.1095098718856988

```

CHECK SENSITIVITY: COVID AFFECTED SEASONS - remove Covid-19 affected seasons from dataset (2019-2020, 2020/2021) - re-train best model and compare to previous result with full dataset to assess whether Covid affected seasons have different driving forces

```

[67]: #Remove Covid-19 shortened season from dataset
      df_reduced = df[df['Season'] != '20202021']
      df_reduced = df_reduced[df_reduced['Season'] != '20192020']
      df_reduced.head()

```

```

[67]:
      Team      Season  GP  W  L  OT  P    P%  RW  ROW  ...  \
62  Tampa Bay Lightning  20182019  82  62  16  4  128  0.780  49  56  ...
63      Calgary Flames  20182019  82  50  25  7  107  0.652  45  50  ...
64      Boston Bruins  20182019  82  49  24  9  107  0.652  38  47  ...
65  Washington Capitals  20182019  82  48  26  8  104  0.634  39  44  ...
66  New York Islanders  20182019  82  48  27  7  103  0.628  37  43  ...

      Win% 3 Goal Game  Wins 1 Goal Game  Wins 2 Goal Game  Wins 3 Goal Game  \
62              0.789              24              8              30
63              0.786              16              12              22
64              0.676              21              5              23
65              0.533              19              13              16
66              0.613              19              10              19

      Loss 1 Goal Game  Loss 2 Goal Game  Loss 3 Goal Game  OT Loss 1 Goal Game  \

```

62	3	5	8	4
63	8	11	6	7
64	6	7	11	9
65	5	7	14	8
66	6	9	12	7

	SAT For/GP	SAT Agst/GP
62	46.658537	43.780488
63	48.073171	41.256098
64	47.158537	41.707317
65	45.292683	47.060976
66	43.231707	47.121951

[5 rows x 80 columns]

```
[68]: #TRAIN BEST MODEL (M5_reduced)

#MODEL M5RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳ df[reg_col_name], test_size=0.3, random_state=2)

M5RR_reduced = Ridge(alpha=0.1)
M5RR_reduced.fit(X_train, y_train)

y_pred = M5RR_reduced.predict(X_test)
```

```
[69]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
print('Adjusted R-Squared:', 1 - (1-metrics.r2_score(y_test,
↳ y_pred))*(len(y_train)-1)/(len(y_train)-X_train.shape[1]-1))
```

Mean Absolute Error: 0.02511585765061583
Mean Squared Error: 0.0011778810039069912
Root Mean Squared Error: 0.03432027103487079
R-Squared: 0.8521876267111418

Adjusted R-Squared: 0.8417362467816265

[]: