

CIND820_project

November 13, 2021

INITIALIZE - import Python modules *****

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn import linear_model
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge, RidgeCV
```

```
[2]: pip install xgboost
```

Collecting xgboost

Using cached xgboost-1.5.0-py3-none-manylinux2014_x86_64.whl (173.5 MB)

Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages
(from xgboost) (1.18.4)

Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages
(from xgboost) (1.4.1)

Installing collected packages: xgboost

Successfully installed xgboost-1.5.0

Note: you may need to restart the kernel to use updated packages.

```
[3]: import xgboost
from xgboost import XGBRegressor
```

IMPORT DATA - import data from csv files - combine into 2 dataframes: -
Regular season - Postseason - check summary statistics for resulting dataframes

```
[4]: #Import data

df_RS_sum = pd.read_csv('Summary_RS.csv', encoding='latin-1')
df_RS_pen = pd.read_csv('Penalties_RS.csv', encoding='latin-1')
df_RS_tgg = pd.read_csv('Team Goal-Games_RS.csv', encoding='latin-1')
df_RS_satc = pd.read_csv('SAT Counts_RS.csv', encoding='latin-1')
df_RS_satp = pd.read_csv('SAT Percentages_RS.csv', encoding='latin-1')
```

```
df_PS_sum = pd.read_csv('Summary_PS.csv', encoding='latin-1')
df_PS_pen = pd.read_csv('Penalties_PS.csv', encoding='latin-1')
df_PS_tgg = pd.read_csv('Team Goal-Games_PS.csv', encoding='latin-1')
df_PS_satc = pd.read_csv('SAT Counts_PS.csv', encoding='latin-1')
df_PS_satp = pd.read_csv('SAT Percentages_PS.csv', encoding='latin-1')
```

```
[5]: #Check import sample
df_RS_sum.head()
```

```
[5]:
```

	Team	Season	GP	W	L	T	OT	P	P%	RW	...	\
0	Vegas Golden Knights	20202021	56	40	14	--	2	82	0.732	30	...	
1	Colorado Avalanche	20202021	56	39	13	--	4	82	0.732	35	...	
2	Carolina Hurricanes	20202021	56	36	12	--	8	80	0.714	27	...	
3	Florida Panthers	20202021	56	37	14	--	5	79	0.705	26	...	
4	Pittsburgh Penguins	20202021	56	37	16	--	3	77	0.688	29	...	

	GA	GF/GP	GA/GP	PP%	PK%	Net PP%	Net PK%	Shots/GP	SA/GP	FOW%
0	122	3.39	2.18	17.8	86.8	16.7	89.6	32.7	27.3	49.5
1	132	3.52	2.36	22.7	83.1	21.3	83.6	34.6	25.4	51.6
2	134	3.13	2.39	25.6	85.2	22.0	89.2	32.0	28.2	53.9
3	151	3.36	2.70	20.5	79.8	18.4	82.1	34.9	30.0	50.2
4	155	3.45	2.77	23.7	77.4	21.1	81.3	30.1	30.0	49.3

[5 rows x 23 columns]

```
[6]: #Create regular season dataframe

#Rename column with shared name but different data
df_RS_satp = df_RS_satp.rename({'GF':'5v5 GF', 'GA':'5v5 GA'}, axis=1)

#Drop duplicate columns
df_RS_pen = df_RS_pen.drop(['GP','W','L','T','OT','P'], axis=1)
df_RS_satc = df_RS_satc.drop(['GP'], axis=1)
df_RS_satp = df_RS_satp.drop(['GP','P','P%'], axis=1)
df_RS_tgg = df_RS_tgg.drop(['GP','W','L','T','OT','P','P%'], axis=1)

#Merge dataframes to create one regular season dataframe
df = pd.merge(df_RS_sum, df_RS_pen, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_satc, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_satp, how='inner', on=['Team','Season'])
df = pd.merge(df, df_RS_tgg, how='inner', on=['Team','Season'])

#Drop irrelevant columns (Ties not applicable in seasons being studied)
df = df.drop(['T'], axis=1)

#Convert season column to string
```

```
df['Season'] = df['Season'].apply(str)

#Convert SAT for/against to rate per game
df['SAT For/GP'] = df['SAT For'] / df['GP']
df['SAT Agst/GP'] = df['SAT Agst'] / df['GP']
```

```
[7]: #Check summary statistics for regular season dataframe
df.describe()
```

```
[7]:
```

	GP	W	L	OT	P	P%	\
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000	
mean	74.311688	37.155844	28.707792	8.448052	82.759740	0.556929	
std	10.380676	8.935594	7.828994	2.906028	17.944379	0.094176	
min	56.000000	15.000000	12.000000	2.000000	37.000000	0.275000	
25%	69.000000	31.000000	24.000000	7.000000	72.000000	0.488000	
50%	82.000000	37.000000	28.000000	8.000000	81.500000	0.569000	
75%	82.000000	44.000000	34.000000	10.000000	97.750000	0.622750	
max	82.000000	62.000000	56.000000	15.000000	128.000000	0.780000	

	RW	ROW	S/O Win	GF	...	Win% 3 Goal Game	\
count	154.000000	154.000000	154.000000	154.000000	...	154.000000	
mean	28.707792	34.298701	2.857143	215.642857	...	0.499370	
std	8.100613	8.792547	1.693620	38.727923	...	0.163777	
min	11.000000	11.000000	0.000000	124.000000	...	0.032000	
25%	23.000000	28.000000	2.000000	189.250000	...	0.385500	
50%	29.000000	34.000000	3.000000	221.000000	...	0.515500	
75%	35.000000	41.000000	4.000000	243.000000	...	0.618000	
max	49.000000	56.000000	9.000000	319.000000	...	0.875000	

	Wins 1 Goal Game	Wins 2 Goal Game	Wins 3 Goal Game	Loss 1 Goal Game	\
count	154.000000	154.000000	154.000000	154.000000	
mean	16.012987	7.733766	13.409091	7.564935	
std	4.216350	2.985898	5.188743	2.976872	
min	5.000000	0.000000	1.000000	1.000000	
25%	13.000000	6.000000	10.000000	5.000000	
50%	16.000000	7.000000	14.000000	8.000000	
75%	19.000000	10.000000	17.000000	9.000000	
max	25.000000	16.000000	30.000000	15.000000	

	Loss 2 Goal Game	Loss 3 Goal Game	OT Loss 1 Goal Game	SAT For/GP	\
count	154.000000	154.000000	154.000000	154.000000	
mean	7.733766	13.409091	8.448052	44.869488	
std	3.031517	5.099806	2.906028	3.125934	
min	2.000000	2.000000	2.000000	37.321429	
25%	5.250000	10.000000	7.000000	42.664634	
50%	8.000000	13.000000	8.000000	44.622570	
75%	10.000000	16.000000	10.000000	46.908537	

max	16.000000	30.000000	15.000000	53.609756
-----	-----------	-----------	-----------	-----------

	SAT Agst/GP
count	154.000000
mean	44.867717
std	3.010061
min	35.053571
25%	42.871951
50%	45.009364
75%	47.088415
max	53.914634

[8 rows x 75 columns]

[8]: *#Create postseason dataframe following same process as regular season*

```
df_PS_satp = df_PS_satp.rename({'GF': '5v5 GF', 'GA': '5v5 GA'}, axis=1)

df_PS_pen = df_PS_pen.drop(['GP', 'W', 'L', 'T', 'OT', 'P'], axis=1)
df_PS_satc = df_PS_satc.drop(['GP'], axis=1)
df_PS_satp = df_PS_satp.drop(['GP', 'P', 'P%'], axis=1)
df_PS_tgg = df_PS_tgg.drop(['GP', 'W', 'L', 'T', 'P', 'P%'], axis=1)

df2 = pd.merge(df_PS_sum, df_PS_pen, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_satc, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_satp, how='inner', on=['Team', 'Season'])
df2 = pd.merge(df2, df_PS_tgg, how='inner', on=['Team', 'Season'])

df2 = df2.drop(['T'], axis=1)

df2['Season'] = df2['Season'].apply(str)

df2['SAT For/GP'] = df2['SAT For'] / df2['GP']
df2['SAT Agst/GP'] = df2['SAT Agst'] / df2['GP']
```

[9]: df2.describe()

[9]:	GP	W	L	P	P%	RW \
count	88.000000	88.000000	88.000000	88.000000	88.000000	88.000000
mean	10.727273	5.363636	5.329545	5.117261	2.505977	4.738636
std	6.629310	4.849889	1.981069	8.387413	3.631763	4.450221
min	3.000000	0.000000	3.000000	0.000000	0.000000	0.000000
25%	5.750000	2.000000	4.000000	0.333000	0.333000	1.000000
50%	8.500000	4.000000	4.000000	0.607500	0.612500	3.000000
75%	15.000000	8.250000	6.250000	6.500000	3.250000	7.000000
max	27.000000	18.000000	12.000000	36.000000	14.000000	16.000000

	ROW	S/O Win	GF	GA	...	5v5 S%+Sv%	\
count	88.000000	88.000000	88.000000	88.000000	...	88.000000	
mean	2.409091	16.272727	29.534091	14.671023	...	98.932955	
std	4.173359	21.981041	18.536047	17.144004	...	3.018781	
min	0.000000	0.000000	4.000000	0.750000	...	90.600000	
25%	0.000000	0.000000	16.000000	2.652500	...	97.150000	
50%	0.000000	7.000000	23.000000	3.355000	...	99.850000	
75%	3.250000	26.250000	42.000000	22.250000	...	101.000000	
max	17.000000	86.000000	77.000000	82.000000	...	103.800000	

	Wins 2 Goal Game	Wins 3 Goal Game	Loss 1 Goal Game	Loss 2 Goal Game	\
count	88.000000	88.000000	88.000000	88.000000	
mean	1.829545	1.397727	2.136364	1.795455	
std	2.354851	1.847893	1.709889	1.576837	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	1.000000	1.000000	
50%	1.000000	1.000000	2.000000	1.500000	
75%	2.000000	2.000000	3.000000	3.000000	
max	12.000000	9.000000	6.000000	7.000000	

	Loss 3 Goal Game	OT Loss 1 Goal Game	Unnamed: 18	SAT For/GP	\
count	88.000000	88.000000	40.000000	88.000000	
mean	1.397727	0.806818	0.075000	47.657206	
std	1.255240	1.239744	0.266747	6.483257	
min	0.000000	0.000000	0.000000	36.571429	
25%	0.000000	0.000000	0.000000	43.221429	
50%	1.000000	0.000000	0.000000	46.900219	
75%	2.000000	1.250000	0.000000	51.723214	
max	5.000000	5.000000	1.000000	69.500000	

	SAT Agst/GP
count	88.000000
mean	47.749411
std	6.219736
min	36.400000
25%	44.250000
50%	47.431373
75%	50.987500
max	66.900000

[8 rows x 72 columns]

INITIAL DATA ASSESSMENT - reduce dataframes to set of variables expected to be of primary interest. Preferentially select variables as rates or ratios from each category of statistics. - create heat map to check correlation between selected variables *****

```
[10]: #Reduce regular season dataframe for initial heatmap
df_reg = df[['W','L','OT','P%', 'RW','ROW','S/O Win','GF/GP','GA/
↳GP','PP%','PK%','Pen Drawn/60','Pen Taken/60','SAT For','SAT_
↳Agst','SAT%','USAT %','Win% 1 Goal Game','Win% 2 Goal Game', 'Win% 3 Goal_
↳Game']]

[11]: #Create correlation heat map for regular season
import matplotlib
matplotlib.use('Agg')
import matplotlib.pyplot as plt
matplotlib.style.use('ggplot')
import seaborn as sns
corr = df_reg.corr()
sns_plot = sns.heatmap(corr, cmap="Blues", annot=True)

fig = sns_plot.get_figure()
fig.savefig("reg_heat.png")

[12]: #Reduce postseason dataframe for initial heatmap
df_post = df2[['W','L','OT','P%', 'RW','ROW','GF/GP','GA/GP','PP%','PK%','Pen_
↳Drawn/60','Pen Taken/60','SAT For','SAT Agst','SAT%','USAT %','Win% 1 Goal_
↳Game','Win% 2 Goal Game', 'Win% 3 Goal Game']]

[13]: #Create correlation heat map for postseason
corr = df_post.corr()
sns_plot = sns.heatmap(corr, cmap="Blues", annot=True)

fig = sns_plot.get_figure()
fig.savefig("post_heat.png")
```

MODELLING: REGULAR SEASON Initial model: - Train initial model for regular season using goal-based input variables - Test multiple types of regression models to find best fit. Considering linear regression, XGBoost regression, ridge regression Additional models: - Add additional variables to initial model in categorical groups to test for improved models.

- Additional categorical groups include shot metrics, penalties, special teams, records in close games - Additional models without using goal metrics

```
[57]: #TRAIN INITIAL MODEL MO
#Linear Regression
#Inputs: goal based variables
#Target: Regular season P%

#Target Variable = Roint Percentage
reg_col_name = 'P%'
```

```

#Select feature variables
feature_names = ['GF/GP', 'GA/GP']

#Use 70/30 train/test split
X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

#Train model & predict test set
MO = linear_model.LinearRegression()
MO.fit(X_train, y_train)

y_pred = MO.predict(X_test)

```

```

[58]: #Check model metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.019301933879415853
 Mean Squared Error: 0.0005326454997485497
 Root Mean Squared Error: 0.02307911392901707
 R-Squared: 0.9331582772976956

```

[14]: #MODEL MO_XGB
#XGBoost Regression
#Inputs: goal based variables
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

MO_XGB = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
    max_depth = 4, n_estimators = 100)

MO_XGB.fit(X_train, y_train)

y_pred = MO_XGB.predict(X_test)

```

```

[15]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))

```

```
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.024832683020449716
 Mean Squared Error: 0.0010186796962558568
 Root Mean Squared Error: 0.031916761995162614
 R-Squared: 0.8721658104466374

```
[16]: #MODEL MORR
      #Ridge Regression
      #Inputs: goal based variables
      #Target: Regular season P%

      reg_col_name = 'P%'

      feature_names = ['GF/GP', 'GA/GP']

      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

      MORR = Ridge(alpha=0.1)
      MORR.fit(X_train, y_train)

      y_pred = MORR.predict(X_test)
```

```
[17]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.019349022473394435
 Mean Squared Error: 0.0005375438606248156
 Root Mean Squared Error: 0.02318499214200461
 R-Squared: 0.93254358163324

```
[ ]: #Continue regular season prediction adding additional variables to goal-based
    ↪model. Linear Regression dropped due to suspected overfitting in M0.
```

```
[18]: #MODEL M1
      #XGBoost Regression
      #Inputs: goal based variables + shot-based metrics
      #Target: Regular season P%

      reg_col_name = 'P%'
```



```

feature_names = ['GF/GP', 'GA/GP', 'SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M1 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
    max_depth = 2, n_estimators = 100)
M1.fit(X_train, y_train)

y_pred = M1.predict(X_test)

```

```

[19]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.024218783079309674
Mean Squared Error: 0.0010224722293729596
Root Mean Squared Error: 0.03197611967348383
R-Squared: 0.8716898851885204

```

[20]: #MODEL M1RR
#Ridge Regression
#Inputs: goal based variables + shot-based metrics
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'SAT For/GP', 'SAT Agst/GP', 'SAT%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M1RR = Ridge(alpha=1)
M1RR.fit(X_train, y_train)

y_pred = M1RR.predict(X_test)

```

```

[21]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.020269714102022034

Mean Squared Error: 0.0006176650868208685
Root Mean Squared Error: 0.024852868784526035
R-Squared: 0.9224891631006638

```
[22]: #MODEL M2
#XGBoost Regression
#Inputs: goal based variables + special teams performance
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'PP%', 'PK%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

M2 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
                    max_depth = 2, n_estimators = 100)
M2.fit(X_train, y_train)

y_pred = M2.predict(X_test)
```

```
[23]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.02436847849602396
Mean Squared Error: 0.0010208861280756123
Root Mean Squared Error: 0.03195130870677463
R-Squared: 0.8718889251562758

```
[24]: #MODEL M3
#XGBoost Regression
#Inputs: goal based variables + penalty rates
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)
```

```

M3 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
    max_depth = 2, n_estimators = 100)
M3.fit(X_train,y_train)

y_pred = M3.predict(X_test)

```

```

[25]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.025027555993262767
 Mean Squared Error: 0.0010153331597157153
 Root Mean Squared Error: 0.031864292863889435
 R-Squared: 0.8725857675617071

```

[26]: #MODEL M4
      #XGBoost Regression
      #Inputs: goal based variables + game-score based records
      #Target: Regular season P%

      reg_col_name = 'P%'

      feature_names = ['GF/GP', 'GA/GP', 'Win% 1 Goal Game', 'Win% 2 Goal Game', 'Win%
    ↪3 Goal Game']

      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
    ↪df[reg_col_name], test_size=0.3, random_state=2)

      M4 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
    ↪learning_rate = 0.1,
    max_depth = 2, n_estimators = 200)
      M4.fit(X_train,y_train)

      y_pred = M4.predict(X_test)

```

```

[27]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.02110979611315626
 Mean Squared Error: 0.0006942845494889962
 Root Mean Squared Error: 0.02634927986661108

R-Squared: 0.912874181129202

```
[28]: #MODEL M4RR
#Ridge Regression
#Inputs: goal based variables + game-score based records
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['GF/GP', 'GA/GP', 'Win% 1 Goal Game', 'Win% 2 Goal Game', 'Win% 3 Goal Game']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names], df[reg_col_name], test_size=0.3, random_state=2)

M4RR = Ridge(alpha=0.1)
M4RR.fit(X_train, y_train)

y_pred = M4RR.predict(X_test)
```

```
[29]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.012573191184253846
Mean Squared Error: 0.00021586526343618027
Root Mean Squared Error: 0.014692353910663201
R-Squared: 0.9729110523106408

```
[ ]: #Continue regular season modeling, without goal based inputs
```

```
[30]: #MODEL M5
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage
#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names], df[reg_col_name], test_size=0.3, random_state=2)

M5 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0, learning_rate = 0.3,
```

```

        max_depth = 10, n_estimators = 100)
M5.fit(X_train,y_train)

y_pred = M5.predict(X_test)

```

```

[31]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.05060893886647325
 Mean Squared Error: 0.0037228397631501094
 Root Mean Squared Error: 0.06101507816228796
 R-Squared: 0.5328205659654223

```

[32]: #MODEL M5RR
      #Ridge Regression
      #Inputs: shot metrics, shooting/save percentage
      #Target: Regular season P%

      reg_col_name = 'P%'

      feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
      ↪df[reg_col_name], test_size=0.3,random_state=2)

      M5RR = Ridge(alpha=0.1)
      M5RR.fit(X_train,y_train)

      y_pred = M5RR.predict(X_test)

```

```

[33]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 0.03219329522690666
 Mean Squared Error: 0.0017621542671799556
 Root Mean Squared Error: 0.041978021239452865
 R-Squared: 0.7788671321899294

```

[34]: #MODEL M6
      #XGBoost Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penatly rates

```

```

#Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳ df[reg_col_name], test_size=0.3, random_state=2)

M6 = XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.5,
↳ learning_rate = 0.1,
                    max_depth = 40, n_estimators = 100)
M6.fit(X_train, y_train)

y_pred = M6.predict(X_test)

```

```

[35]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

```

Mean Absolute Error: 0.04500114227355795
Mean Squared Error: 0.0030528095528431706
Root Mean Squared Error: 0.05525223572710131
R-Squared: 0.6169027060391592

```

```

[48]: #MODEL M6RR
      #Ridge Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
      #Target: Regular season P%

reg_col_name = 'P%'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
↳ df[reg_col_name], test_size=0.3, random_state=2)

M6RR = Ridge(alpha=0.1)
M6RR.fit(X_train, y_train)

y_pred = M6RR.predict(X_test)

```

```
[49]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.025264003393739975
Mean Squared Error: 0.0011659666653443162
Root Mean Squared Error: 0.03414625404556576
R-Squared: 0.8536827579283653

POSTSEASON MODELS - Model postseason, using wins as target variable. - Use input variables matching regular season models: initial base model and other most successful regular season models
- Results to be compared with regular season models to assess impact of small postseason sample sizes *****

```
[54]: #MODEL P6
      #XGBoost Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
      #Target: Postseason W

      reg_col_name = 'W'

      feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

      X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
      ↪df2[reg_col_name], test_size=0.3, random_state=2)

      P6 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.0,
      ↪learning_rate = 0.1,
                        max_depth = 4, n_estimators = 100)
      P6.fit(X_train, y_train)

      y_pred = P6.predict(X_test)
```

```
[55]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 2.89310469561153
Mean Squared Error: 14.81850461444183
Root Mean Squared Error: 3.8494810837880253
R-Squared: 0.3156790913513181

```
[56]: #MODEL P6RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%']

X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
↳df2[reg_col_name], test_size=0.3,random_state=2)

P6 = Ridge(alpha=0.1)
P6.fit(X_train,y_train)

y_pred = P6.predict(X_test)
```

```
[57]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

```
Mean Absolute Error: 2.8631163201994414
Mean Squared Error: 12.537485594919337
Root Mean Squared Error: 3.540831201133335
R-Squared: 0.42101691380361106
```

```
[58]: #MODEL P0
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
↳df2[reg_col_name], test_size=0.3,random_state=2)

P0 = XGBRegressor(objective ='reg:squarederror', colsample_bytree = 1.0,
↳learning_rate = 0.1,
max_depth = 4, n_estimators = 100)
P0.fit(X_train,y_train)

y_pred = P0.predict(X_test)
```



```
[59]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 3.720087496218858
Mean Squared Error: 20.83882289016658
Root Mean Squared Error: 4.564955957089463
R-Squared: 0.03765983232412051

```
[66]: #MODEL PORR
      #Ridge Regression
      #Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
      #Target: Postseason W

      reg_col_name = 'W'

      feature_names = ['GF/GP', 'GA/GP']

      X_train, X_test, y_train, y_test = train_test_split(df2.loc[:, feature_names],
      ↪df2[reg_col_name], test_size=0.3, random_state=2)

      PORR = Ridge(alpha=0.5)
      PORR.fit(X_train, y_train)

      y_pred = PORR.predict(X_test)
```

```
[67]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 3.76161208247753
Mean Squared Error: 23.979697311732068
Root Mean Squared Error: 4.896906912708477
R-Squared: -0.10738624985763812

PREDICT POSTSEASON SUCCESS USING REGULAR SEASON INPUT VARIABLES - Use regular season input variables to predict postseason success. Chose variables based on results of previous models. *****

```
[69]: #Use regular season statistics for input variables
```

```

RS = df[['Team', 'Season', 'RW', 'P', 'GF/GP', 'GA/GP', 'Win% 1 Goal Game',
        ↪ 'Win% 2 Goal Game', 'Win% 3 Goal Game', 'SAT For/GP', 'SAT Agst/GP', 'SAT%',
        ↪ '5v5 S%', '5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]

#Use postseason W as target variable
PS = df2[['Team', 'Season', 'P%', 'W']]

#Create new dataframe
df3 = pd.merge(RS, PS, on=['Team', 'Season'])

```

```

[72]: #MODEL RP6
#XGBoost Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
        ↪ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
        ↪ df3[reg_col_name], test_size=0.3, random_state=2)

RP6 = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.5,
        ↪ learning_rate = 0.1,
                    max_depth = 10, n_estimators = 100)
RP6.fit(X_train, y_train)

y_pred = RP6.predict(X_test)

```

```

[73]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
        ↪ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

```

Mean Absolute Error: 4.9501098659303455
Mean Squared Error: 36.16998315436581
Root Mean Squared Error: 6.014148581001787
R-Squared: -0.32916209897835835

```

```

[74]: #MODEL RP6RR
#Ridge Regression
#Inputs: shot metrics, shooting/save percentage, special teams, penatly rates
#Target: Postseason W

reg_col_name = 'W'

```

```

feature_names = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%',
↳ 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
↳ df3[reg_col_name], test_size=0.3, random_state=2)

RP6RR = Ridge(alpha=0.01)
RP6RR.fit(X_train, y_train)

y_pred = RP6RR.predict(X_test)

```

```

[75]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

```

Mean Absolute Error: 5.125942556086683
Mean Squared Error: 39.63840663970098
Root Mean Squared Error: 6.295903957312324
R-Squared: -0.45661853212733194

```

```

[76]: #MODEL RPO
#XGBoost Regression
#Inputs: goal rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
↳ df3[reg_col_name], test_size=0.3, random_state=2)

RPO = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.5,
↳ learning_rate = 0.1,
                    max_depth = 10, n_estimators = 100)
RPO.fit(X_train, y_train)

y_pred = RPO.predict(X_test)

```

```

[77]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↳ y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))

```

Mean Absolute Error: 5.422038421862656
Mean Squared Error: 48.07547664870346
Root Mean Squared Error: 6.933648148608599
R-Squared: -0.7666610785817534

```
[78]: #MODEL RPORR
#Ridge Regression
#Inputs: goal rates
#Target: Postseason W

reg_col_name = 'W'

feature_names = ['GF/GP', 'GA/GP']

X_train, X_test, y_train, y_test = train_test_split(df3.loc[:, feature_names],
    ↪df3[reg_col_name], test_size=0.3, random_state=2)

RPO = Ridge(alpha=0.1)
RPO.fit(X_train, y_train)

y_pred = RPO.predict(X_test)
```

```
[79]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 4.693376720415959
Mean Squared Error: 35.838794498087864
Root Mean Squared Error: 5.986551135510985
R-Squared: -0.3169916921618132

CHECK REGULAR SEASON TO POSTSEASON CORRELATIONS - check correlation of each variable used as input variable to previous model between regular season and postseason. - higher correlation may suggest higher predictive power in postseason prediction models.

```
[80]: #Check correlation between key variables: regular season vs postseason

#Combine data
reg_shot = df[['Team', 'Season', 'SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%',
    ↪'5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]
post_shot = df2[['Team', 'Season', 'SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5
    ↪S%', '5v5 Sv%', 'PP%', 'PK%', 'Pen Drawn/60', 'Pen Taken/60']]
```

```

post_shot = post_shot.rename(columns={'SAT For/GP': 'SAT For/GP P', 'SAT Agst/
↳GP': 'SAT Agst/GP P', 'SAT%': 'SAT% P', '5v5 S%': '5v5 S% P', '5v5 Sv%': '5v5
↳Sv% P', 'PP%': 'PP% P', 'PK%': 'PK% P', 'Pen Drawn/60': 'Pen Drawn/60 P', 'Pen
↳Taken/60': 'Pen Taken/60 P'})
compare = pd.merge(reg_shot, post_shot, on=['Team', 'Season'])

#Check correlations
from scipy.stats import pearsonr
stats = ['SAT For/GP', 'SAT Agst/GP', 'SAT%', '5v5 S%', '5v5 Sv%', 'PP%', 'PK%',
↳'Pen Drawn/60', 'Pen Taken/60']

for i in stats:
    list1 = compare[i]
    list2 = compare[i + ' P']

    corr, _ = pearsonr(list1, list2)
    print('Pearsons correlation (' + i + '): %.3f' % corr)

```

```

Pearsons correlation (SAT For/GP): 0.261
Pearsons correlation (SAT Agst/GP): 0.254
Pearsons correlation (SAT%): 0.466
Pearsons correlation (5v5 S%): -0.080
Pearsons correlation (5v5 Sv%): 0.043
Pearsons correlation (PP%): -0.063
Pearsons correlation (PK%): -0.106
Pearsons correlation (Pen Drawn/60): 0.393
Pearsons correlation (Pen Taken/60): -0.125

```

CHECK SENSITIVITY: COVID AFFECTED SEASONS - remove Covid-19 affected seasons from dataset (2019-2020, 2020/2021) - re-train base model and compare to previous result with full dataset to assess whether Covid affected seasons have different driving forces

```

[81]: #Remove Covid-19 shortened season from dataset
df_reduced = df[df['Season'] != '20202021']
df_reduced = df_reduced[df_reduced['Season'] != '20192020']
df_reduced.head()

```

```

[81]:
      Team      Season  GP  W  L  OT  P   P%  RW  ROW  ...  \
62 Tampa Bay Lightning 20182019  82  62  16   4  128  0.780  49  56  ...
63   Calgary Flames 20182019  82  50  25   7  107  0.652  45  50  ...
64   Boston Bruins 20182019  82  49  24   9  107  0.652  38  47  ...
65 Washington Capitals 20182019  82  48  26   8  104  0.634  39  44  ...
66 New York Islanders 20182019  82  48  27   7  103  0.628  37  43  ...

```

```

Win% 3 Goal Game  Wins 1 Goal Game  Wins 2 Goal Game  Wins 3 Goal Game  \

```

62	0.789	24	8	30
63	0.786	16	12	22
64	0.676	21	5	23
65	0.533	19	13	16
66	0.613	19	10	19

	Loss 1 Goal Game	Loss 2 Goal Game	Loss 3 Goal Game	OT Loss 1 Goal Game \
62	3	5	8	4
63	8	11	6	7
64	6	7	11	9
65	5	7	14	8
66	6	9	12	7

	SAT For/GP	SAT Agst/GP
62	46.658537	43.780488
63	48.073171	41.256098
64	47.158537	41.707317
65	45.292683	47.060976
66	43.231707	47.121951

[5 rows x 80 columns]

```
[82]: #TRAIN INITIAL MODEL (MO_reduced): GOAL-BASED INDEPENDENT VARIABLES / REGULAR_
      ↪SEASON

      #Target Variable = Roint Percentage
      reg_col_name = 'P%'

      #Use only goal-based variables
      feature_names = ['GF/GP', 'GA/GP']

      #Use 70/30 train/test split
      X_train, X_test, y_train, y_test = train_test_split(df.loc[:, feature_names],
      ↪df[reg_col_name], test_size=0.3, random_state=2)

      #Train model & predict test set
      MO_reduced = XGBRegressor(objective='reg:squarederror', colsample_bytree = 1.
      ↪0, learning_rate = 0.1,
      max_depth = 4, n_estimators = 100)
      MO_reduced.fit(X_train, y_train)

      y_pred = MO_reduced.predict(X_test)
```

```
[83]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
      print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
```

```
print('R-Squared:', metrics.r2_score(y_test, y_pred))
```

Mean Absolute Error: 0.024832683020449716

Mean Squared Error: 0.0010186796962558568

Root Mean Squared Error: 0.031916761995162614

R-Squared: 0.8721658104466374

[]: